

Practical Machine Learning Project

PQ

31 January 2016

1.Executive Summary

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. In this project, we will be using the data collected from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways. The data for this project come from this source:

<http://groupware.les.inf.puc-rio.br/har>.

The goal is to predict the manner in which they did the exercise from a set of 20 records (validate_DF). This is the "classe" variable (values are A, B, C, D, E) in the training set.

2.Loading the Libraries and Getting the Data

```
#Load Libraries
library(caret)

## Warning: package 'caret' was built under R version 3.2.3
## Loading required package: lattice
## Loading required package: ggplot2
library(randomForest)
## Warning: package 'randomForest' was built under R version 3.2.3
## randomForest 4.6-12
## Type rfNews() to see new features/changes/bug fixes.
set.seed(123)

#Load the data required
validate_DF = read.csv("pml-testing.csv", na.strings=c("NA", "#DIV/0!"))
all_data = read.csv("pml-training.csv", na.strings=c("NA", "#DIV/0!"))
```

A quick look at the data. We can see that there quite a number of variables with a lot of missing (NA) values or error values (#DIV/0)

```
summary(all_data)
```

```
summary(all_data$classe)
```

3.Cleaning the Data

Now, we will proceed to tidy up the data before creating our models

```
#Removing variables with high number of NAs
all_data_rownum <- nrow(all_data)
all_data_colnum <- length(all_data)

#Create a new data frame that will be used for the training model
all_data_clean <- all_data

#Remvoving columns/variables with high number (90%) of NA values.
#We are left with 60 columns.
all_data_clean <- all_data_clean[, colSums(is.na(all_data_clean)) <
                                     all_data_rownum * 0.9]

#remove column 1 (index number of the dataset)
all_data_clean<- all_data_clean[,c(-1)]

#Applying the same to the validation data set, using what is the column names
that is left
remaining_col <- names(all_data_clean)
#This dataset does not have the "classe" data
remaining_col <- remaining_col[remaining_col!="classe"]
validate_DF_clean <- validate_DF[remaining_col]
```

4.Partioning the data into 2 sets.

The training data for the model is 60% and testing data set is 40%

```
inTrain = createDataPartition(all_data_clean$classe, p = 0.60)[[1]]
training = all_data_clean[ inTrain,]
testing = all_data_clean[-inTrain,]
```

5.Model Building

The random forest method is used in this model

```
Model1Control <- trainControl(method="cv", number=3, verboseIter=F)
Model1_RF <- train(classe ~ ., data=training, method="rf",
trControl=Model1Control)

Model1_RF$finalModel

##
## Call:
```

```
## randomForest(x = x, y = y, mtry = param$mtry)
##           Type of random forest: classification
##           Number of trees: 500
## No. of variables tried at each split: 41
##
##           OOB estimate of  error rate: 0.12%
## Confusion matrix:
##      A      B      C      D      E  class.error
## A 3348      0      0      0      0 0.0000000000
## B      2 2277      0      0      0 0.0008775779
## C      0      2 2049      3      0 0.0024342746
## D      0      0      4 1925      1 0.0025906736
## E      0      0      0      2 2163 0.0009237875
```

Let's now fit the model with the testing data set, as to ensure that we are not overfitting the model based on the training set.

```
pred_Model1_RF <- predict(Model1_RF, newdata=testing)
```

6.Accuracy of the Mode

The accuracy of the model is 99.94%, which is a very good number. In this case, we will proceed with using this model to predict the "classe" values for the 20 records in the "validate_DF_clean" data set.

```
confusionMatrix(testing$classe,pred_Model1_RF)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction      A      B      C      D      E
##           A 2232      0      0      0      0
##           B      0 1518      0      0      0
##           C      0      2 1365      1      0
##           D      0      0      0 1286      0
##           E      0      0      0      0 1442
##
## Overall Statistics
##
##           Accuracy : 0.9996
##           95% CI : (0.9989, 0.9999)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9995
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
```

##	Class: A	Class: B	Class: C	Class: D	Class: E
## Sensitivity	1.0000	0.9987	1.0000	0.9992	1.0000
## Specificity	1.0000	1.0000	0.9995	1.0000	1.0000
## Pos Pred Value	1.0000	1.0000	0.9978	1.0000	1.0000
## Neg Pred Value	1.0000	0.9997	1.0000	0.9998	1.0000
## Prevalence	0.2845	0.1937	0.1740	0.1640	0.1838
## Detection Rate	0.2845	0.1935	0.1740	0.1639	0.1838
## Detection Prevalence	0.2845	0.1935	0.1744	0.1639	0.1838
## Balanced Accuracy	1.0000	0.9993	0.9998	0.9996	1.0000

7.Prediction / Results

Now, getting the predictions of "classe" based on the random forest model created - whether they are "A","B","C","D" or "E".

```
pred1 <- predict(Model1_RF,validate_DF_clean,type="raw")
pred1

## [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```